



THE UNIVERSITY OF  
MELBOURNE

# COMP90050 Advanced Database Systems

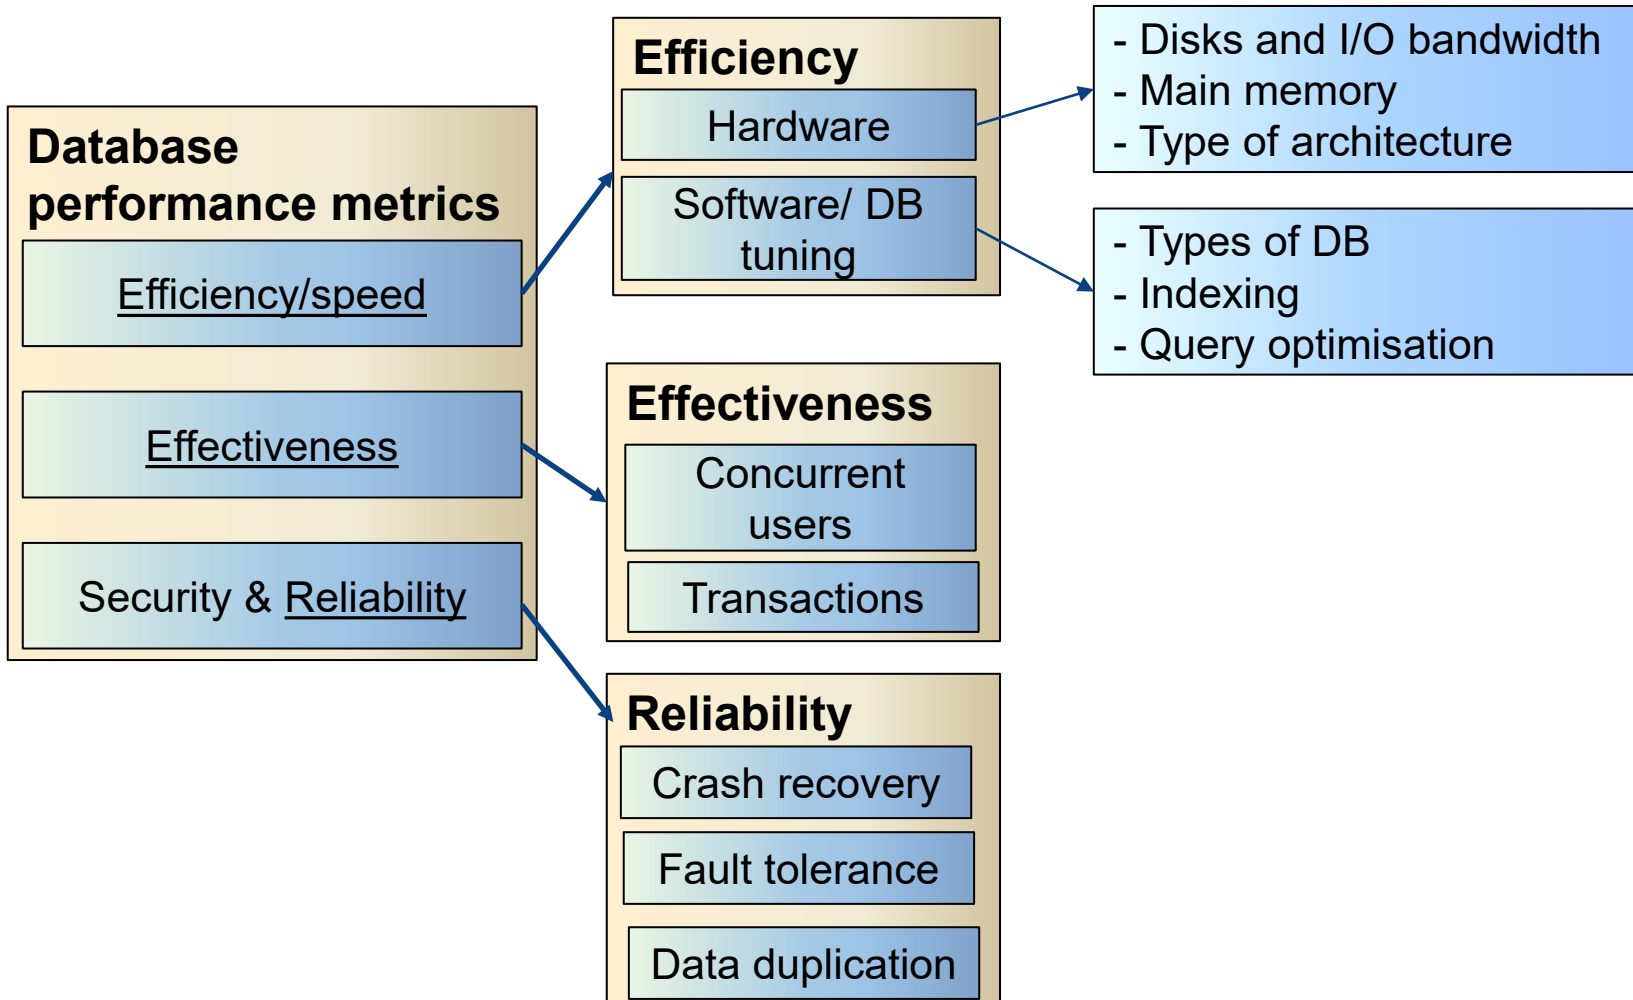
## Winter Semester, 2023

Lecturer: Farhana Choudhury (PhD)

Week 2 part 4-5



# Core Concepts of Database management system





# What needs to be efficient

- DBMS must support:
  - insert/delete/modify record
  - read a particular record (specified using record id)
  - scan all records (possibly with some conditions on the records to be retrieved), or scan a range of records



# A key choice to make

- **DBMS admin** generally **creates indices** to allow almost direct access to individual items
- These are also good during join operations  
if there is a join condition that restricts the number of items to be joined in a table

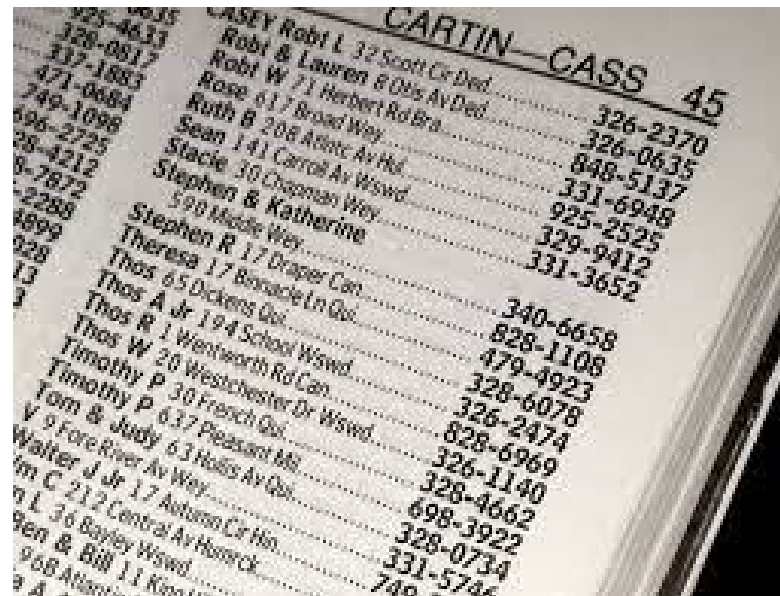


# Topics

- What is indexing
- Different types of indexes
- Usage of indexes
- Indexing in practice

# Indexing is Critical for Efficiency

- Indexing mechanisms used to speed up access to desired data in a similar way to look up a phone book or dictionary
- **Search Key** - attribute or set of attributes used to look up records/rows in a system like an ID of a person






# Indexing is Critical for Efficiency

- An **index file** consists of records (called **index entries**) of the form **search-key, pointer to where data is**
- Index files are typically much smaller than the original data files and many parts of it are already in memory

# What becomes faster?

□ Disk access becomes faster through:

- records with a specified value in the attribute accessed with minimal disk accesses
- or records with an attribute value falling in a specified range of values can be retrieved with a single seek and then consecutive sequential reads







# Indexing is Critical for Efficiency

- ❑ Insertion time to index is also important
- ❑ Deletion time is important as well
- ❑ No big index rearrangement after insertion and deletion
- ❑ Space overhead needs to be considered for the index itself
- ❑ No single indexing technique is the best. Rather, each technique is best suited to particular applications.



# Indexing is Critical for Efficiency


- Two basic kinds of indices based on search keys:
  - **Ordered indices:** search keys are stored in some order
  - **Hash indices:** search keys are distributed hopefully uniformly across “buckets” using a “function”

# Ordered Index

The records in the indexed file may themselves be stored in some sorted order

- **Clustering index/ Primary index:** in a sequentially ordered file, the index whose search key specifies the sequential order of the file
  - The search key of a primary index is usually but not necessarily the primary key


10101	Srinivasan	Comp. Sci.	65000	
12121	Wu	Finance	90000	
15151	Mozart	Music	40000	
22222	Einstein	Physics	95000	
32343	El Said	History	60000	



# Ordered Index

The records in the indexed file may themselves be stored in some sorted order

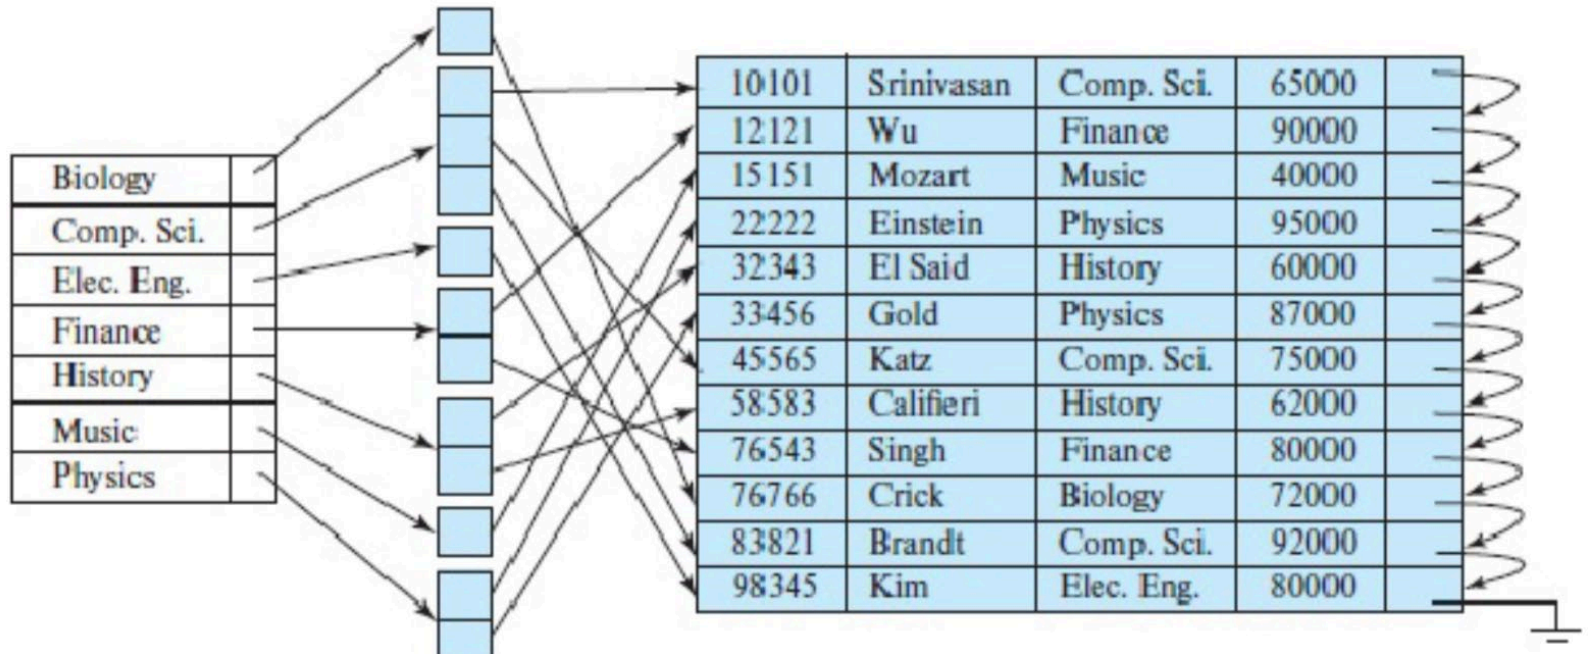
- **Clustering index/ Primary index:** in a sequentially ordered file, the index whose search key specifies the sequential order of the file



# Ordered Index

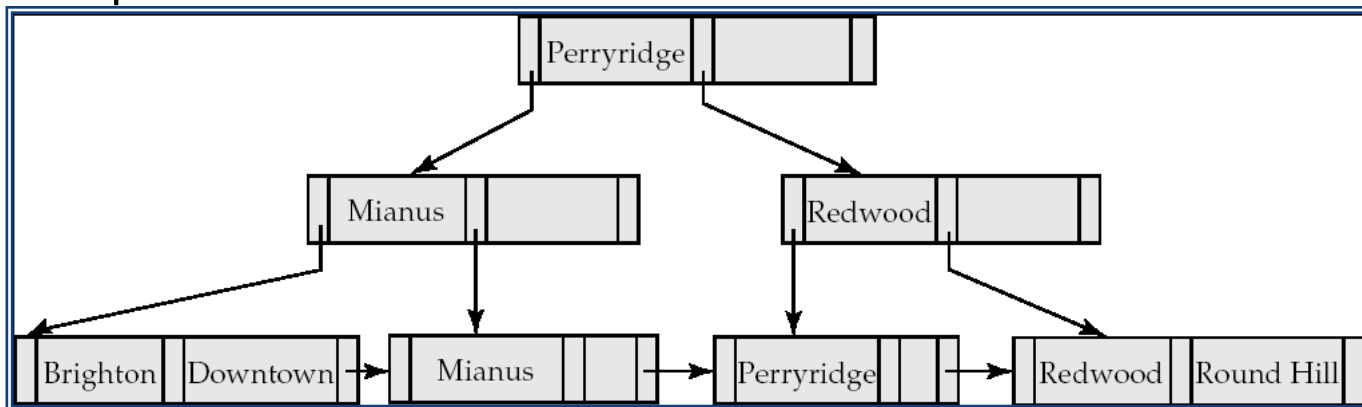
- **Non-clustering index/ Secondary index:** an index whose search key specifies an order different from the sequential order of the file

Secondary indices improve the performance of queries that use keys other than the search key of the clustering index.



# The most popular in DBMS: B+trees

- Why need them:
  - Keeping files in order for fast search ultimately degrades as file grows, since many overflow blocks get created.
  - So binary search on ordered files cannot be done.
  - Periodic reorganization of entire file is required to achieve this.
- Advantage of B<sup>+</sup>-tree index files:
  - automatically reorganizes itself with small, local, changes, in the face of insertions and deletions.
  - Reorganization of entire file is not required to maintain performance.






# B+trees Contd

- Similar to Binary tree in many aspects but the fan out is much higher
- Disadvantage of B<sup>+</sup>-trees:
  - Extra insertion and deletion overhead and space overhead
- Advantages of B<sup>+</sup>-trees outweigh disadvantages for DBMSs
  - B<sup>+</sup>-trees are used extensively

# B+tree structure

**B+ tree of order 2:** each (internal) node has between 2 and  $2^2$  DATA entries (except possibly the root)



**B+ tree of order d:** each (internal) node has between d and  $2^d$  entries (except possibly the root)



Video link: <https://youtu.be/CYKRMz8yzVU>



# How is a B+tree defined?

- It is similar to a binary tree in concept but with a fan out that is defined through a number  $n$
- All paths from root to leaf are of the same length (depth)
- Each node that is not a root or a leaf has between  $\lceil n/2 \rceil$  and  $n$  children
- A leaf node has between  $\lceil (n-1)/2 \rceil$  and  $n-1$  values
- Special cases:
  - If the root is not a leaf, it has at least 2 children.
  - If the root is a leaf (that is, there are no other nodes in the tree), it can have between 0 and  $(n-1)$  values.

# A Single Node

- Typical node




- $K_i$  are the search-key values
- $P_i$  are pointers to children (for non-leaf nodes) or pointers to records or buckets of records (for leaf nodes)
- The search-keys in a node are ordered

$$K_1 < K_2 < K_3 < \dots < K_{n-1}$$

NOTE: Most of the higher level nodes of a B+tree would be in main memory already!

# Example B+tree



B<sup>+</sup>-tree for *account* file ( $n = 5$ )

- Leaf nodes must have between 2 and 4 values ( $\lceil (n-1)/2 \rceil$  and  $n-1$ , with  $n = 5$ ).
- Non-leaf nodes other than root must have between 3 and 5 children ( $\lceil n/2 \rceil$  and  $n$  with  $n = 5$ ).
- Root must have at least 2 children.

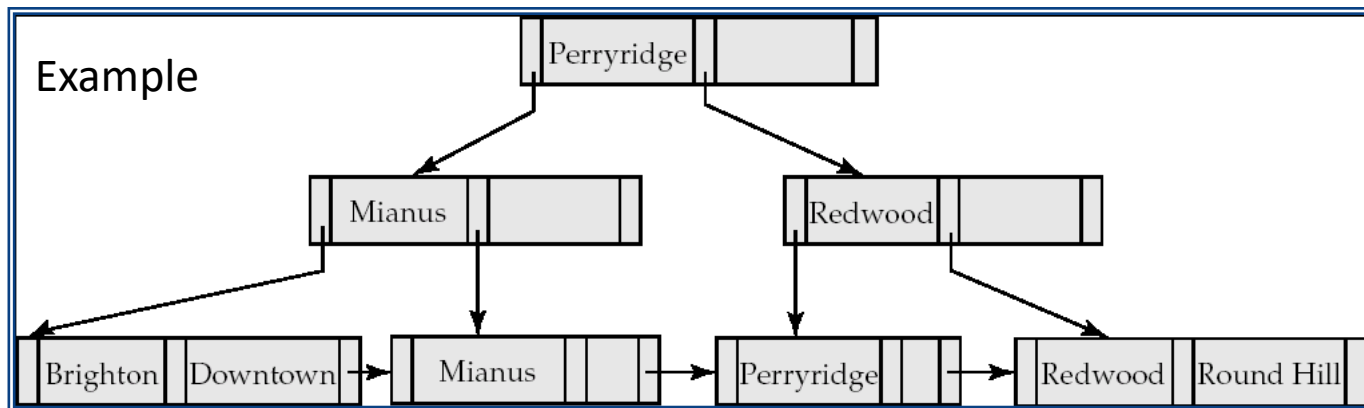
# How to run a query fast then?

□ Finding all records with a search-key value of  $k$ .

1.  $N = \text{root initially}$
2. Repeat
  1. Examine  $N$  for the smallest search-key value  $> k$ .
  2. If such a value exists, assume it is  $K_i$ . Then set  $N = P_i$
  3. Otherwise  $k \geq K_{n-1}$ . Set  $N = P_n$ . Follow pointer.

Until  $N$  is a leaf node

3. If for some  $i$ , key  $K_i = k$  follow pointer  $P_i$  to the desired record or bucket.
4. Else no record with search-key value  $k$  exists.



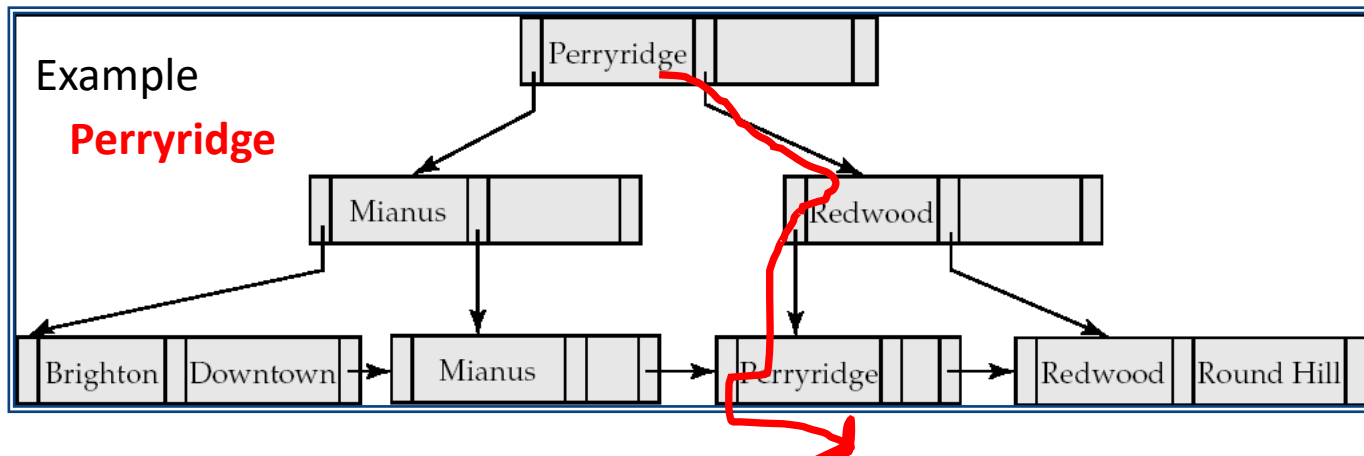
# Recall: running a query on B+trees?

□ Finding all records with a search-key value of  $k$ .

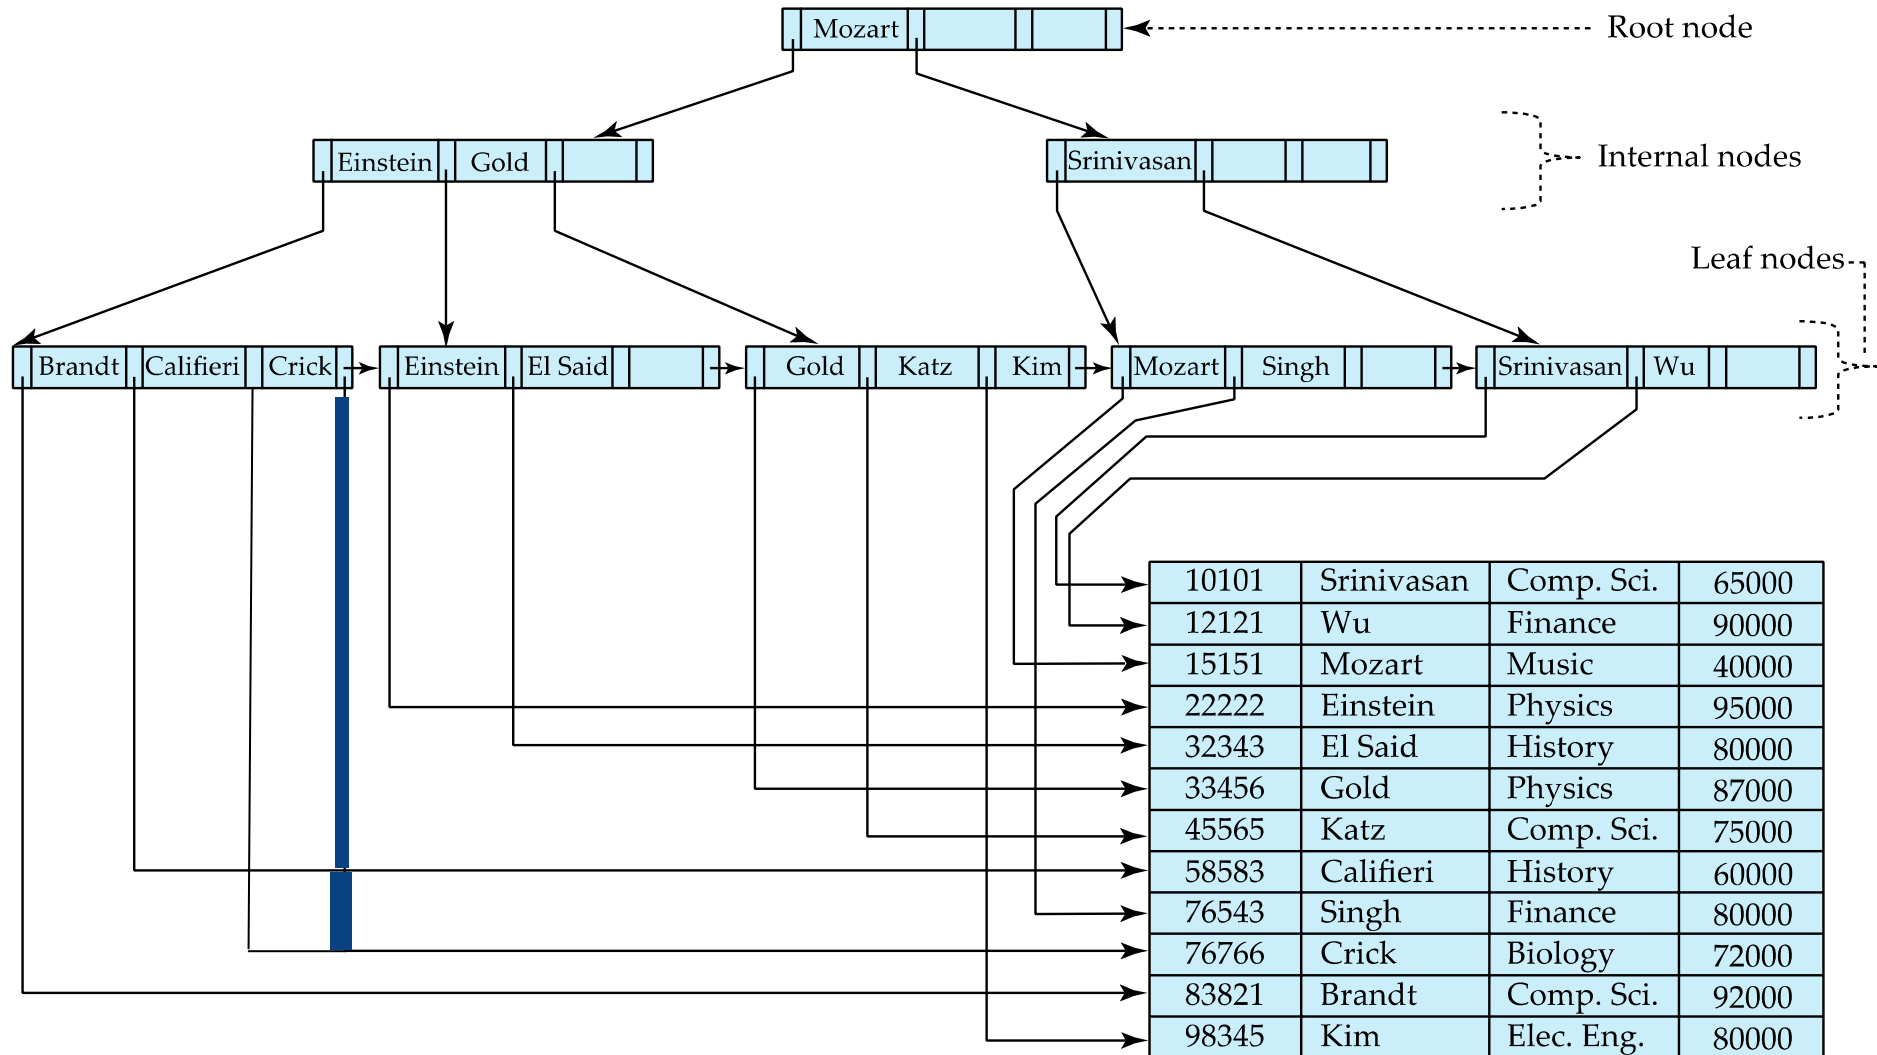
1.  $N = \text{root initially}$
2. Repeat
  1. Examine  $N$  for the smallest search-key value  $> k$ .
  2. If such a value exists, assume it is  $K_i$ . Then set  $N = P_i$
  3. Otherwise  $k \geq K_{n-1}$ . Set  $N = P_n$ . Follow pointer.

Until  $N$  is a leaf node

3. If for some  $i$ , key  $K_i = k$  follow pointer  $P_i$  to the desired record or bucket.
4. Else no record with search-key value  $k$  exists.



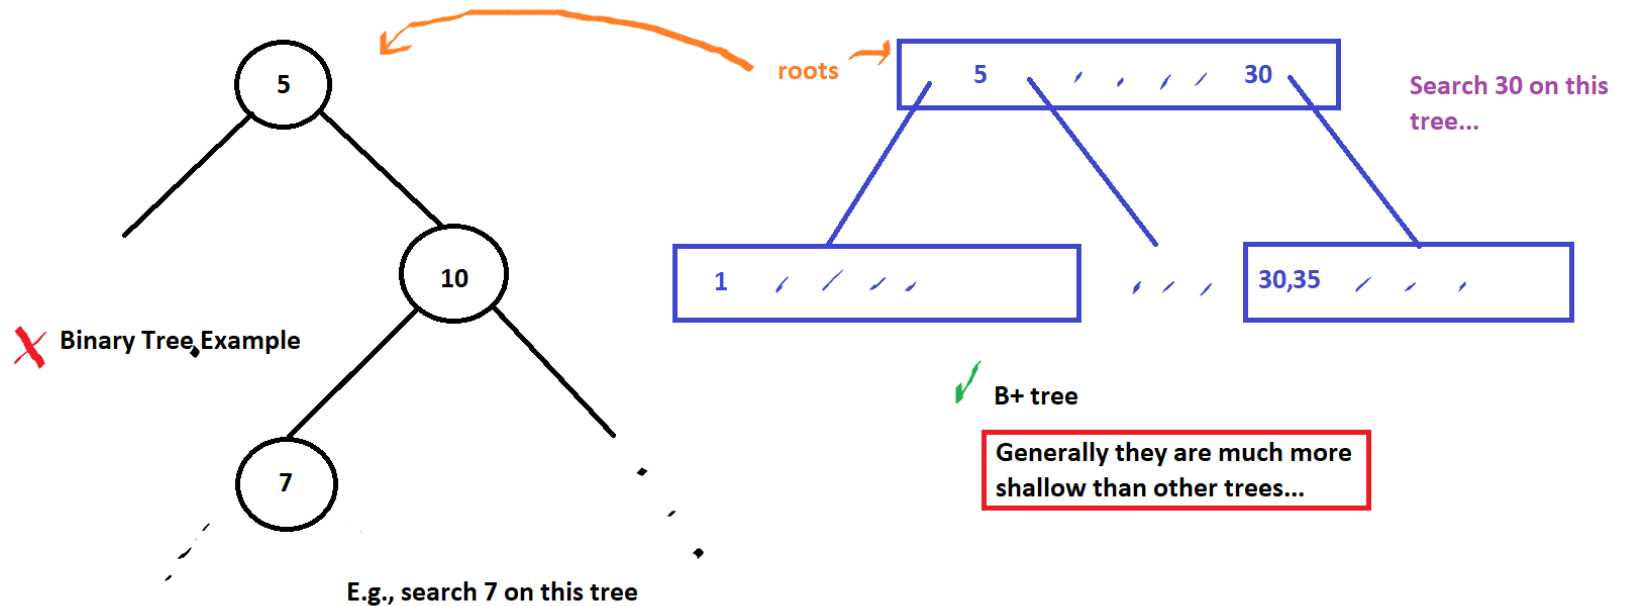
# A full view of a B<sup>+</sup>-Tree



# Metrics

- If there are  $K$  search-key values in the file, the height of the tree is no more than  $\lceil \log_{\lceil n/2 \rceil}(K) \rceil$  and **it would be balanced**
- A node is generally the same size as a disk block, typically 4 kilobytes
  - and  **$n$  is typically around 100** (40 bytes per index entry).
- With 1 million search key values and  $n = 100$ 
  - $\log_{50}(1,000,000) = \underline{\mathbf{4 \text{ nodes are accessed in a lookup}}}$ .
- Contrast this with a balanced binary tree with 1 million search key values — around **20 nodes are accessed in a lookup**
  - above difference is significant since every node access may need a disk I/O, costing around 20 milliseconds

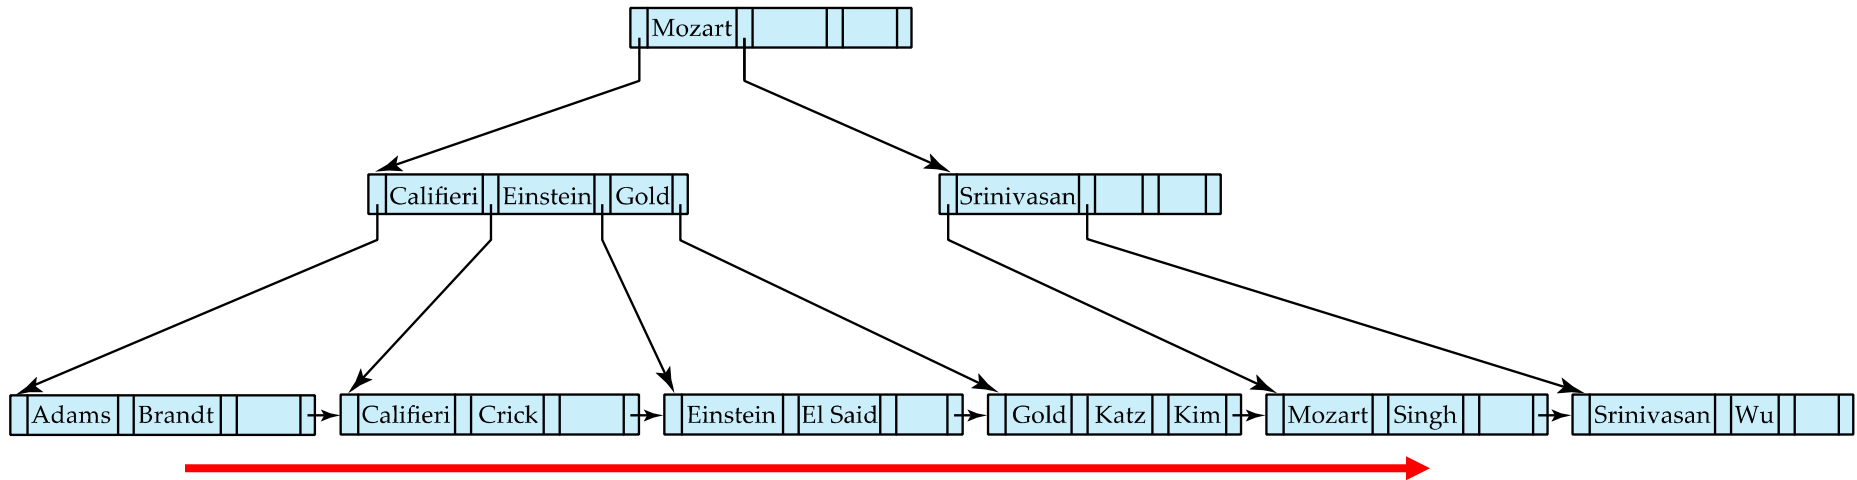
# A Quick Visual Comparison



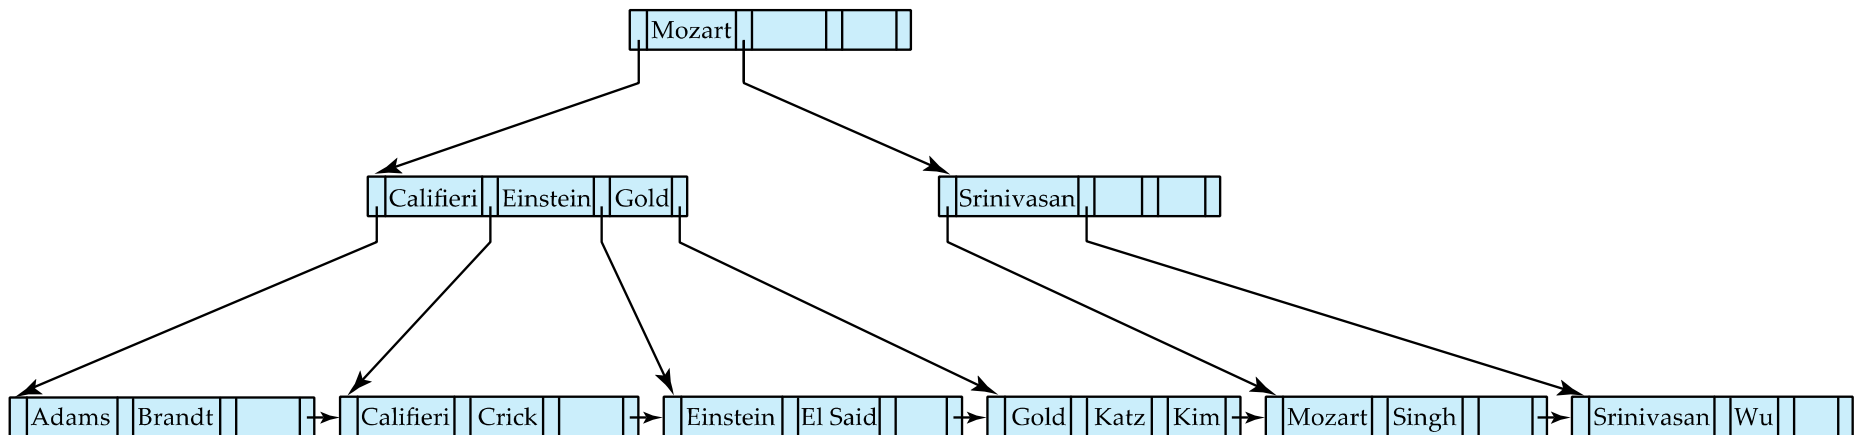
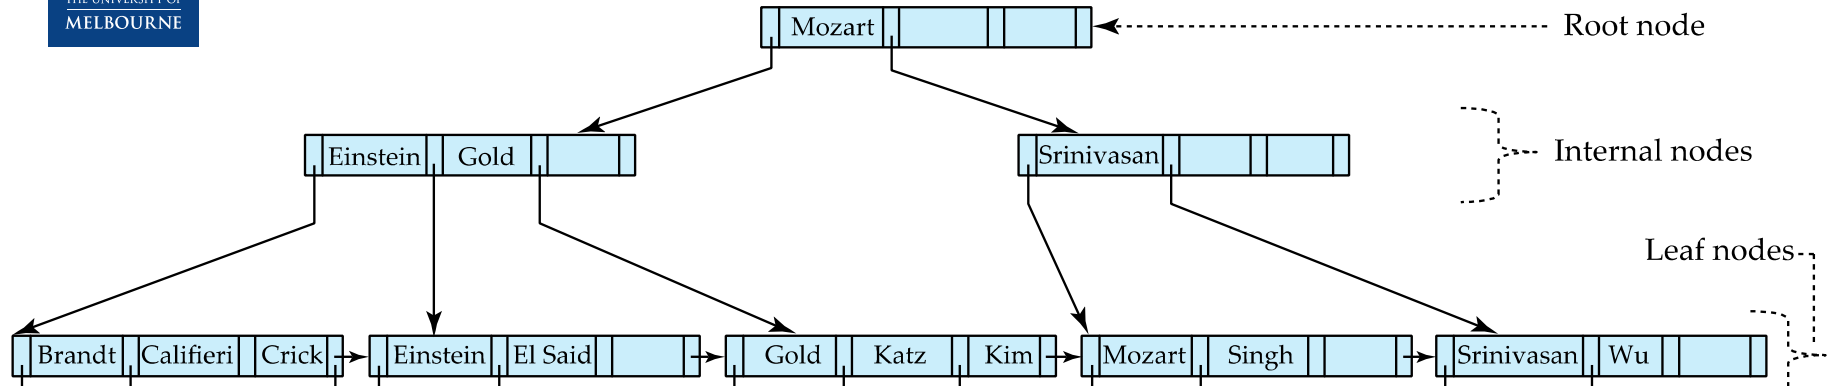


# Range Queries on B<sup>+</sup>-Trees

**Range queries** find all records with search key values in a given range

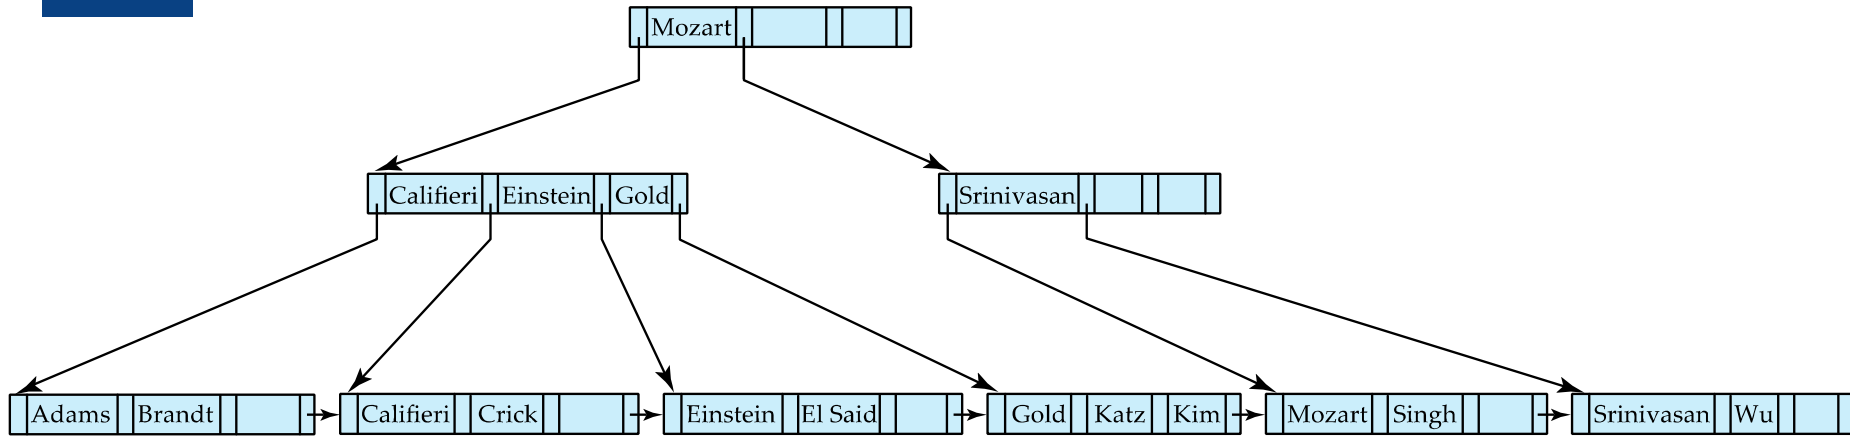


# B<sup>+</sup>-Tree Insertion

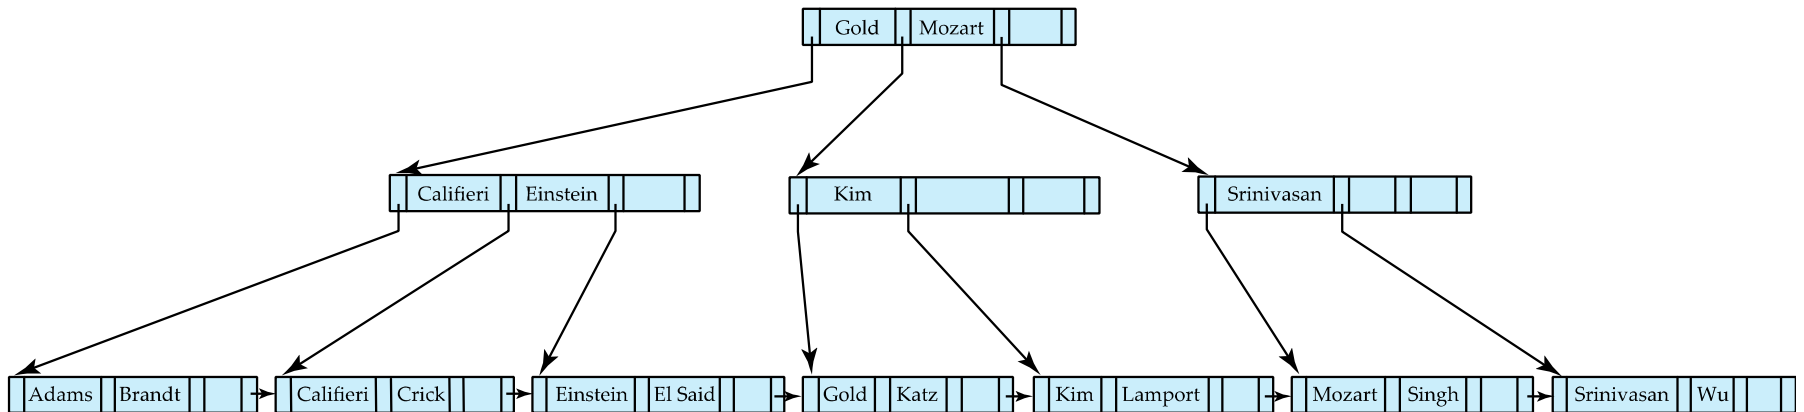


B<sup>+</sup>-Tree before and after insertion of "Adams"

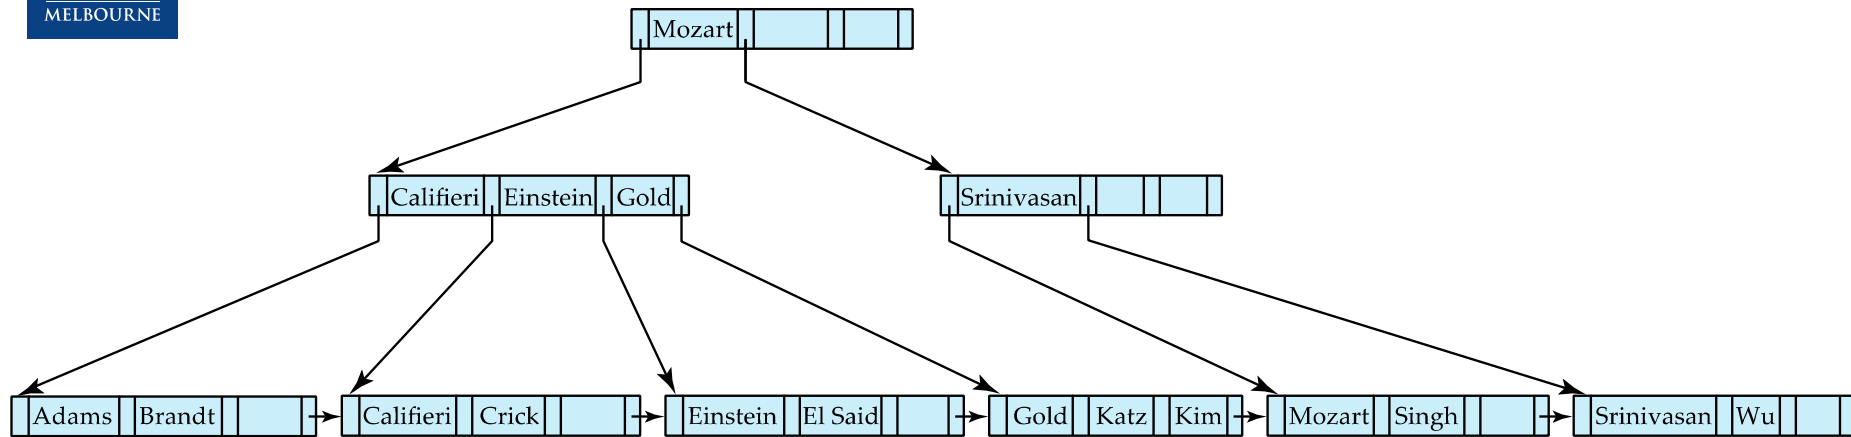
# B<sup>+</sup>-Tree Insertion



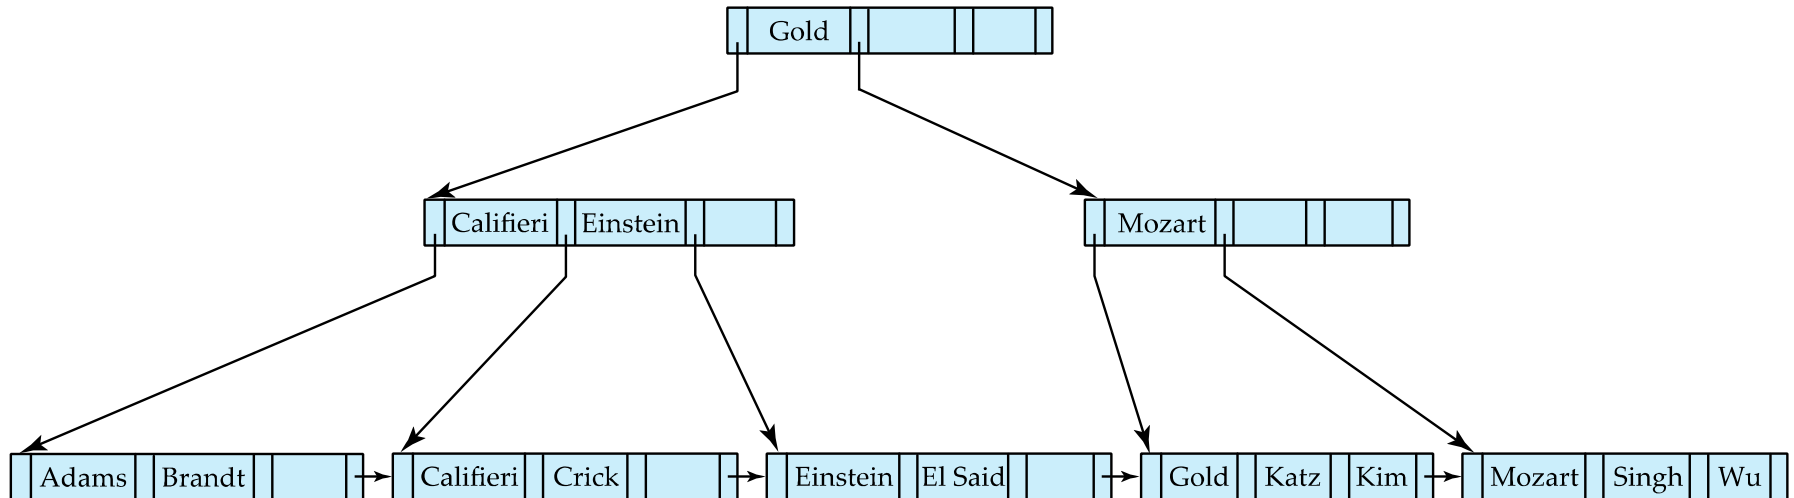
## B<sup>+</sup>-Tree before and after insertion of “Lamport”



# Examples of B<sup>+</sup>-Tree Deletion

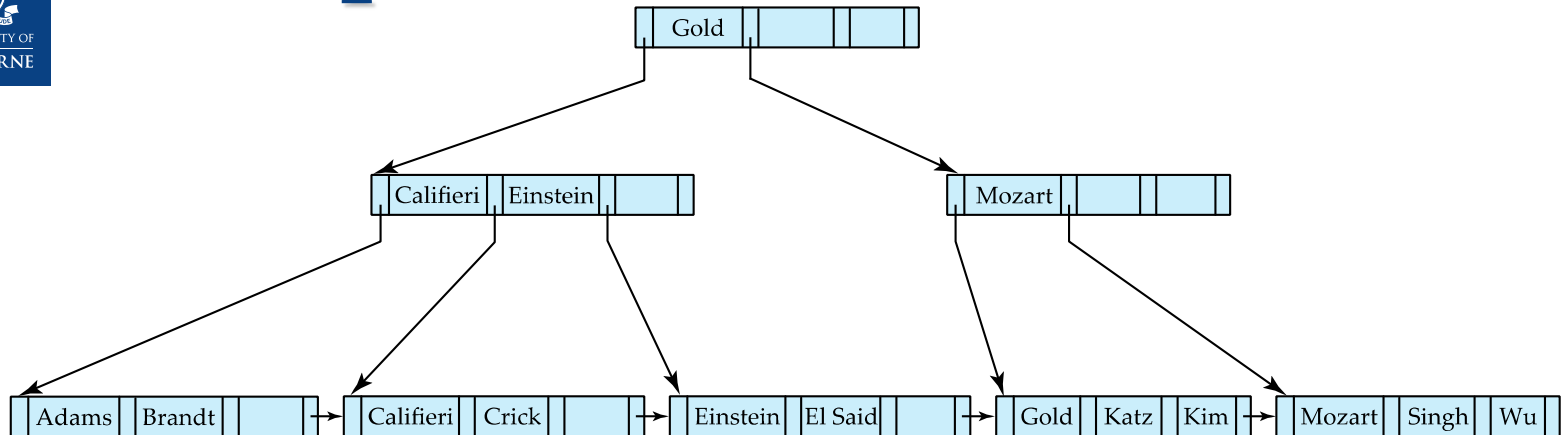


Before and after deleting “Srinivasan”




Deleting “Srinivasan” causes **merging** of under-full leaves

# Examples of B<sup>+</sup>-Tree Deletion




Before and after deleting “Singh” and “Wu”



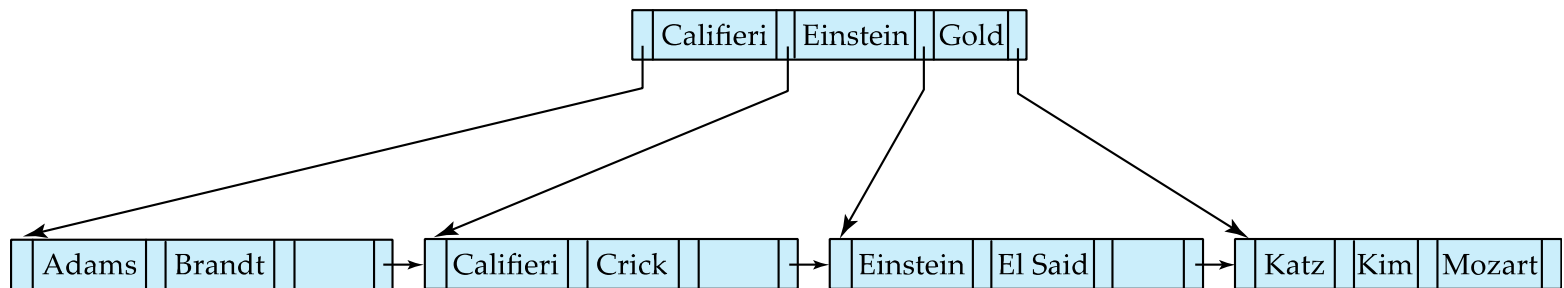
Leaf containing Singh and Wu became underfull, and **borrowed a value** Kim from its left sibling

Search-key value in the parent changes as a result

# Example of B<sup>+</sup>-tree Deletion



**Before and after deletion of “Gold”**



Node with Gold and Katz became underfull, and was merged with its sibling

Parent node becomes underfull, and is merged with its sibling

- Value separating two nodes (at the parent) is pulled down when merging

Root node then has only one child, and is deleted



# B<sup>+</sup>-Tree File Organization

## B<sup>+</sup>-Tree File Organization:

- Leaf nodes in a B<sup>+</sup>-tree file organization store records, instead of pointers to children
- Helps keep data records clustered (ordered) even when there are insertions/deletions/updates
- Insertion and deletion of records are handled in the same way as insertion and deletion of entries in a B<sup>+</sup>-tree index.