



THE UNIVERSITY OF
MELBOURNE

COMP90050 Advanced Database Systems

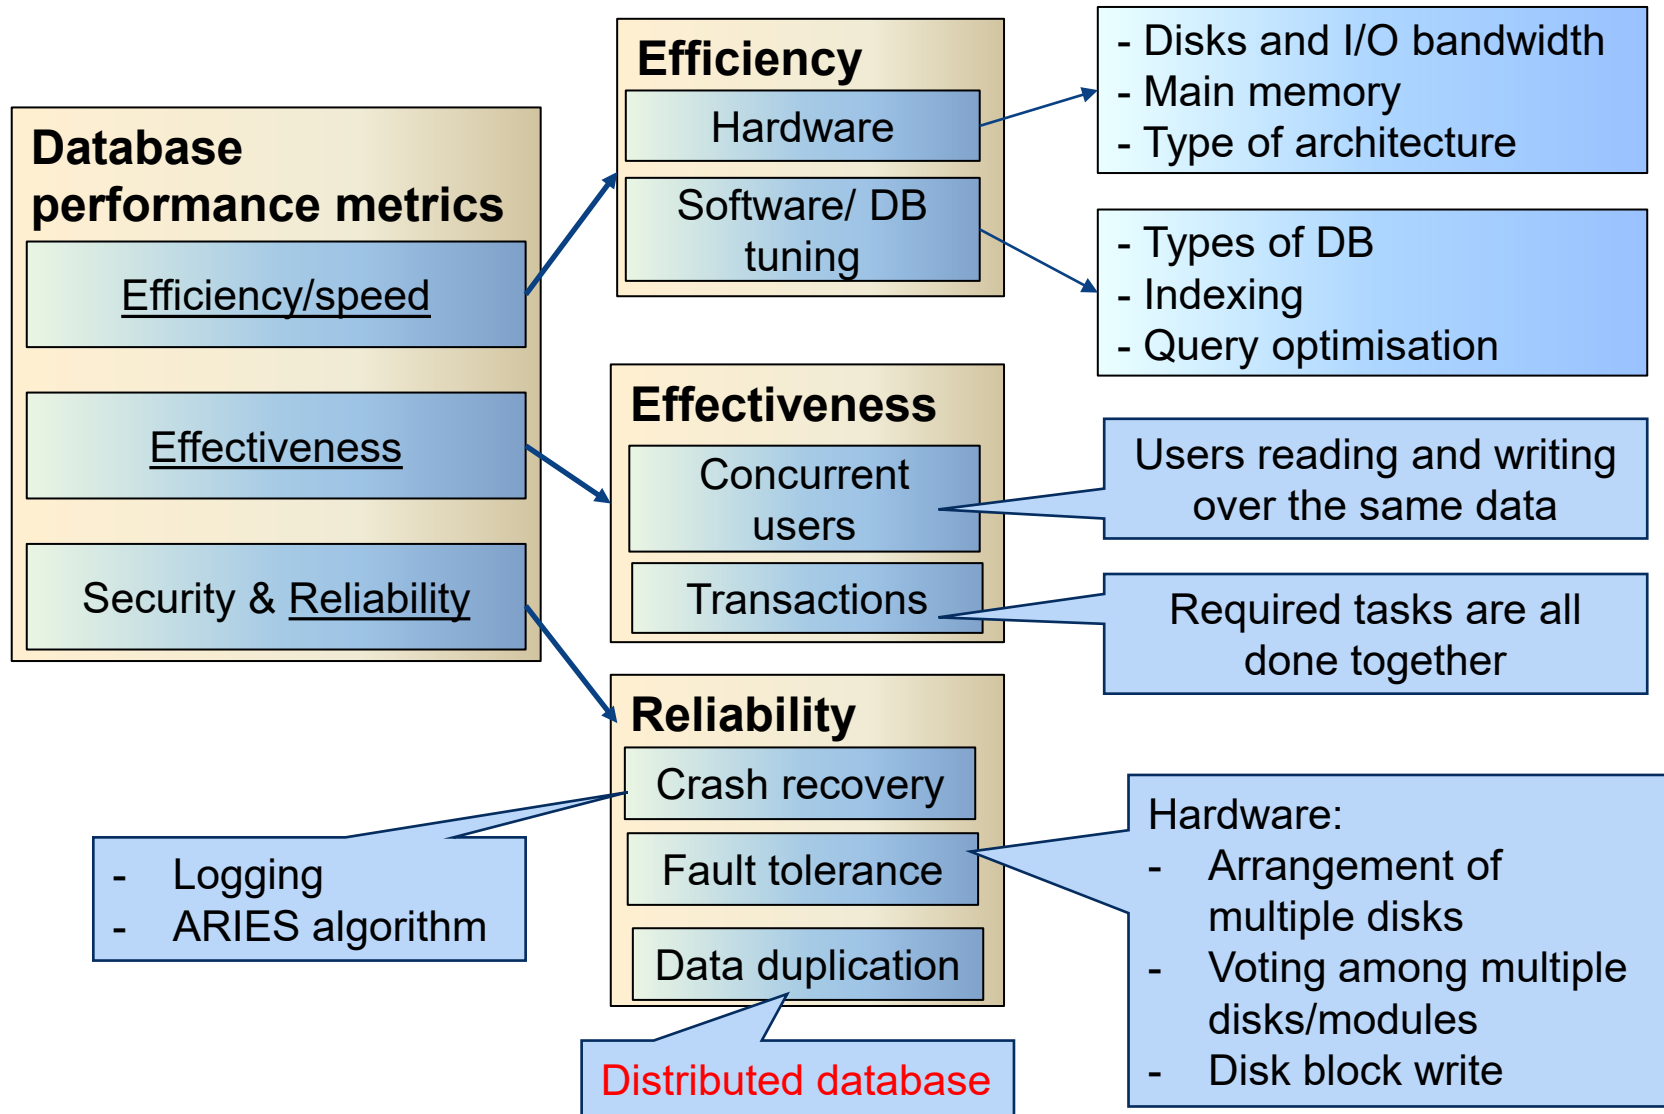
Winter Semester, 2023

Lecturer: Farhana Choudhury (PhD)

Week 4 part 4



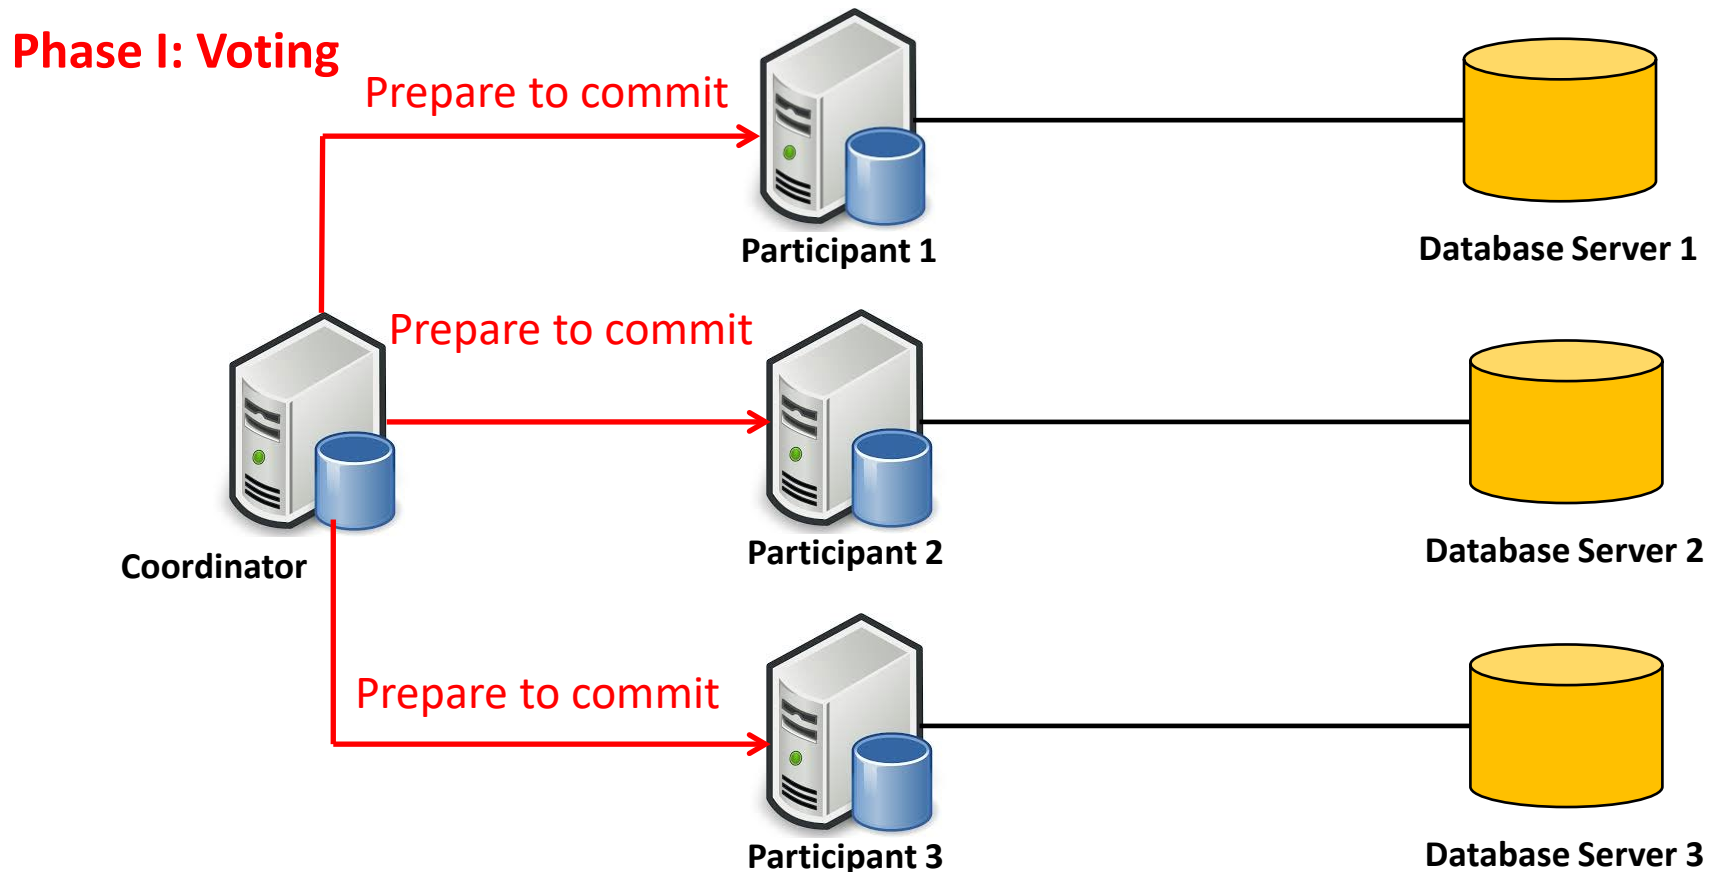
Core Concepts of Database management system



A related concept on distributed databases first - atomicity

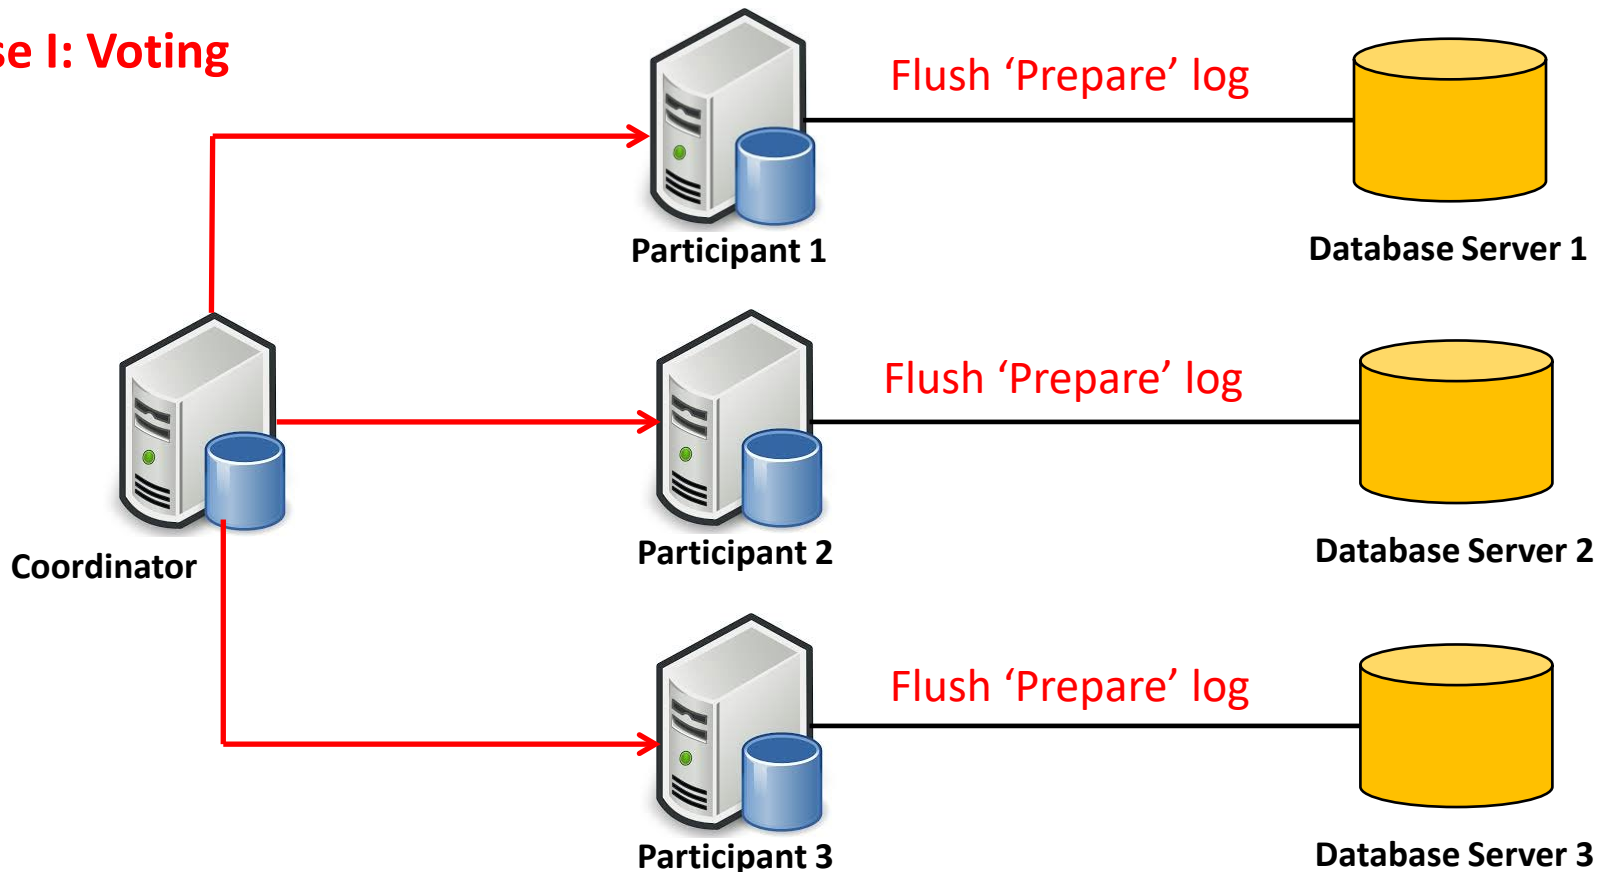
Atomicity in distributed transaction processing

The two-phase commit protocol (2PC) can help achieve atomicity in distributed transaction processing




Two phase commit protocol

Phase I: Voting




Two phase commit protocol

Phase I: Voting



Two phase commit protocol

Phase II: Commit





Two phase commit protocol

- Coordinator or participant can abort transaction
 - If a participant abort, it must inform coordinator
 - If a participant does not respond within a timeout period, coordinator will abort
- If abort, coordinator asks all participants to rollback
- If abort, abort logs are forced to disk at coordinator and all participants

Another related concept on distributed databases - concurrency control



Concurrency control in distributed DBs

- Each server is responsible for applying concurrency control to its own objects
- The members of a collection of servers of distributed transactions are jointly responsible for ensuring that they are performed in a serially equivalent manner
- **BUT** servers independently acting would not work
- If transaction **T is before transaction U** in their conflicting access to objects at one of the servers then:
 - They **must be in that order at all of the servers** whose objects are accessed in a conflicting manner by both T and U
- The central **Coordinator** should assure this



Other Considerations for Locking-based systems

- A local lock manager cannot release any locks until it knows that the transaction has been committed or aborted at all the servers involved in the transaction.
- The objects remain locked and are unavailable for other transactions during the commit protocol.
 - An aborted transaction releases its locks after phase 1 of the protocol.



Timestamp ordering concurrency control revisited

- The coordinator accessed by a transaction issues a **globally unique timestamp**
- The **timestamp is passed with each object access**
- The servers are jointly responsible for ensuring serial equivalence:
 - that is **if T access an object before U, then T is before U at all objects**



Optimistic concurrency control revisited

For distributed transactions to work:

- 1) **Validation takes place in phase 1 of 2PC protocol at each server**
- 2) **Transactions use a globally unique order for validation**

What if objects in different servers are replicas for increased availability

Client 1:	Client 2:	Server
setBalance_B(x,1)		
setBalance _A (y,2)		
	getBalance _A (y) → 2	
	getBalance_A(x) → 0	

Initial balance of x and y is \$0

The **behaviour above cannot occur if A and B did not exist** (that is, if we had only one server)

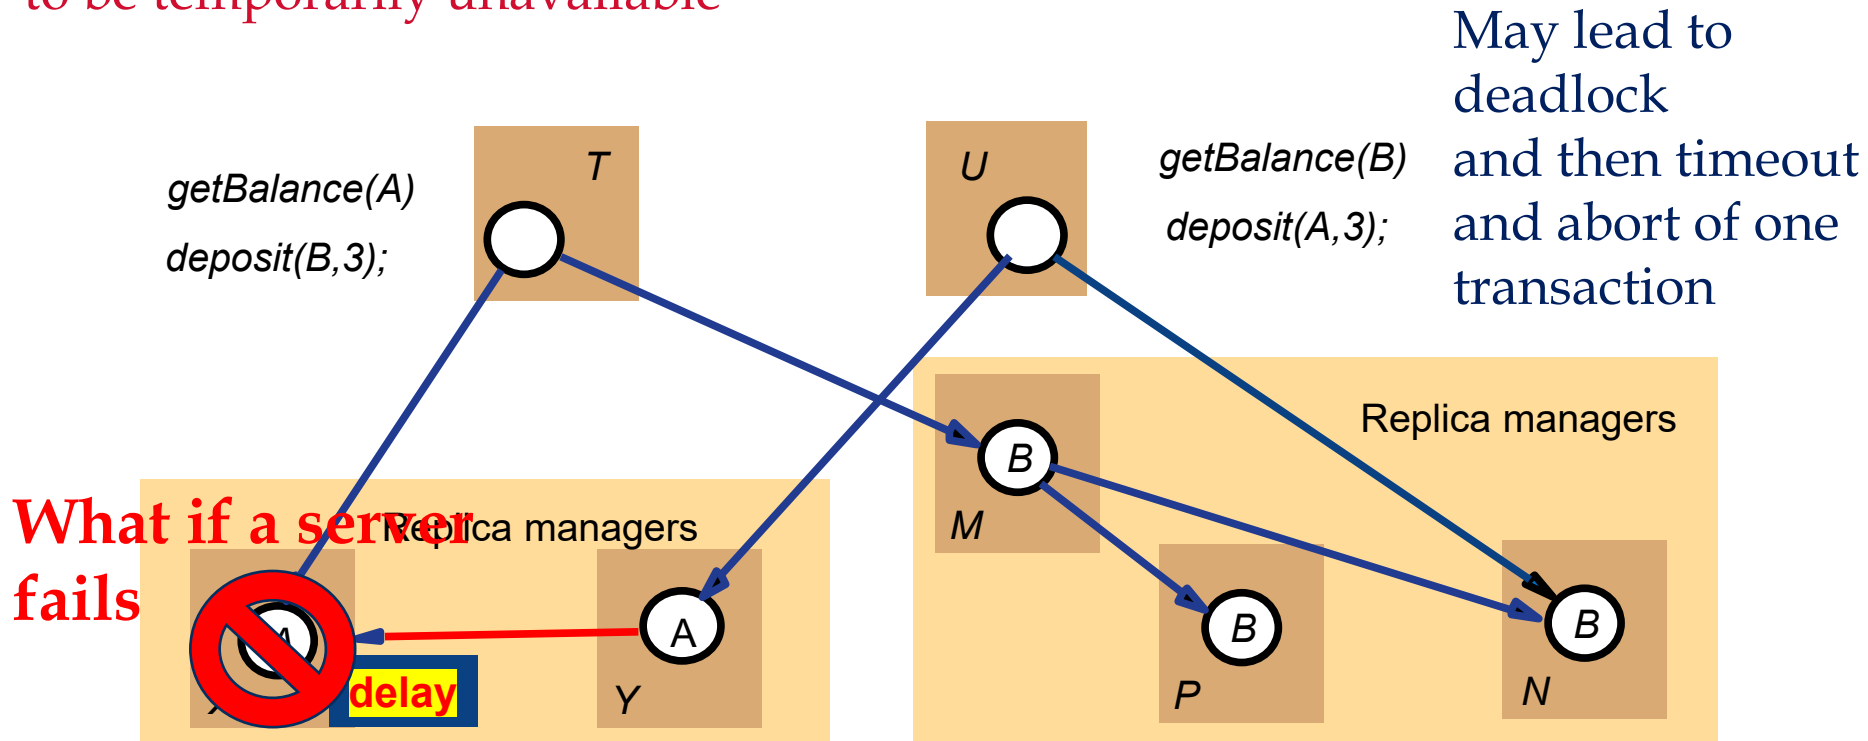


Transactions with replicated data

- the effect of transactions on replicated objects should be the **same as if they had been performed one at a time on a single set of objects**
- this property is called ***one-copy serializability***
- **If all servers are available then no issue – but what if some servers are not available?**

Lets build the solution step by step

The **available copies replication** scheme is designed to allow some servers to be temporarily unavailable

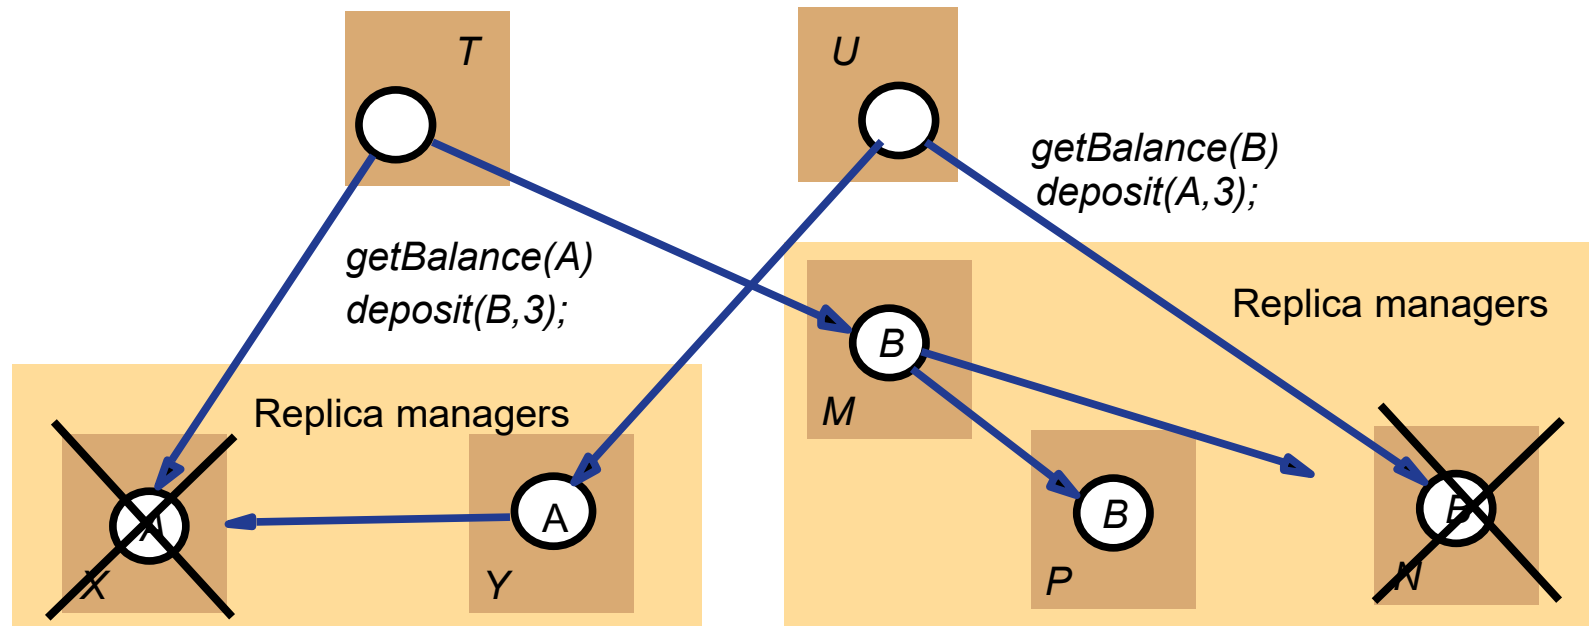


At *X*, *T* has read *A* and has locked it. Therefore *U*'s *deposit* is delayed until *T* finishes. Normally, this leads to good concurrency control only if the servers do not fail....

Cannot the other servers simply be used?

assume that X fails just after T has performed *getBalance*

and N failing just after U has performed *getBalance*



therefore T 's deposit will be performed at M and P (all available)


and U 's deposit will be performed at Y (all available)

NOT GOOD!!

Available copies replication rule

Before a transaction commits, it checks for failed and available servers it has contacted, the set should not change during execution:

- E.g., T would check if X is still available among others.
- We said X fails before T 's *deposit*, in which case, T would have to abort.
- Thus no harm can come from this execution now.



Core Concepts of Database management system

